

# BPF on Computational Storage

Niclas Hedam | [nhed@itu.dk](mailto:nhed@itu.dk)



IT UNIVERSITY OF COPENHAGEN

# Outline

- Introduction (problem, approach, contributions)
- Background
- Architecture
- Mechanisms
- Use-cases
- Conclusion

# Introduction

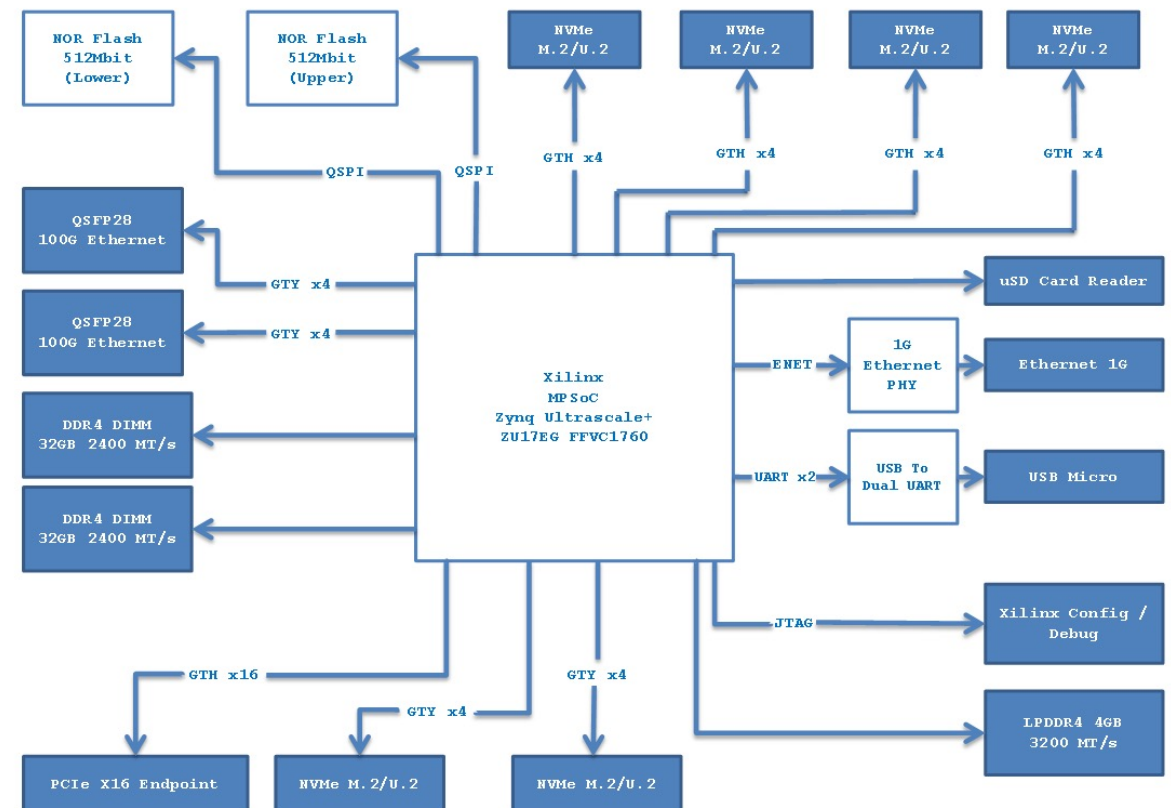
- We are using BPF to orchestrate function calls on computational storage devices (CSDs).
- Our approach is experimental, contributing to the required mechanisms, in the context of use-cases provided by the DAPHNE project.
- So far our contributions include:
  - Architecture for BPF on CSD
  - Extension of the uBPF virtual machine to support registered functions on CSD
  - XDMA mechanisms and BPF execution in Eid-Hermes (QEMU)
  - DAPHNE use-cases

# Background: BPF

- BPF is a assembly like format bytecode
  - Support in clang (and gcc) for C code
    - Not so much for C++
  - **Vendor neutral ISA**
    - Possible execution at various steps of I/O path: kernel, NIC, CSD
    - **But:**
      - Verification in kernel, NIC or CSD are different, e.g., focus on confidentiality for in-kernel execution vs. focus on integrity for CSD verification, bytecode vs. C-level analysis
        - Verification on CSD is an open-problem!
      - No ARM64 JIT for BPF
        - Neither in-kernel or in uBPF!
      - Different VMs in kernel, NIC or CSD.
        - uBPF for user-space VM on CSD.
          - Who is in charge of uBPF? Pull request from Eideticom blocked for weeks...
          - Universal uBPF vs. CSD-specific BPF VM?

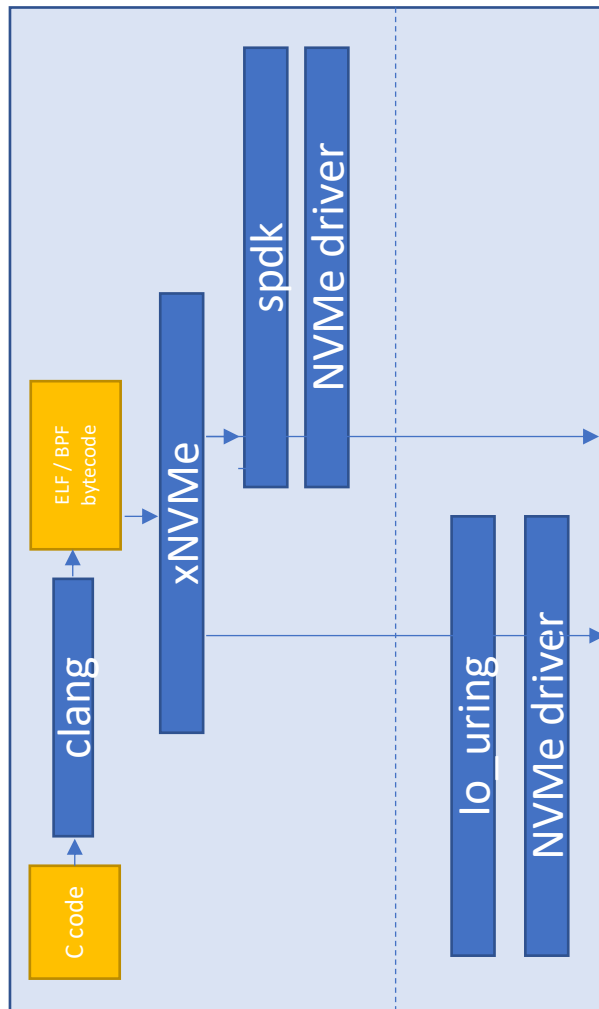
# Background: CSD

- We will experiment on the Daisy OpenSSD.
- Specs
  - 1.5 GHz Quad-core ARM main processor
  - 0.6 GHz Dual-core ARM real time processor
  - 6 NVMe expansion slots




# Architecture

- Programs are loaded into a BPF virtual machine on CSD.
  - We use a modified version of the uBPF VM by iovisor.
- VMs are managed by a *loader*, which is in charge of loading programs into the VM and managing the state and interconnect.
  - The loader is also responsible for exposing *registered functions*.
  - We are building this loader ourselves at ITU.



Host


NVM  
admin  
data



SSD


Registers external functions in BPF VM  
Allocates/deallocates memory  
Maintains state

NVM  
admin  
data  
CS




Daisy CSD

NVM  
admin  
data



SSD

NVM  
admin  
data



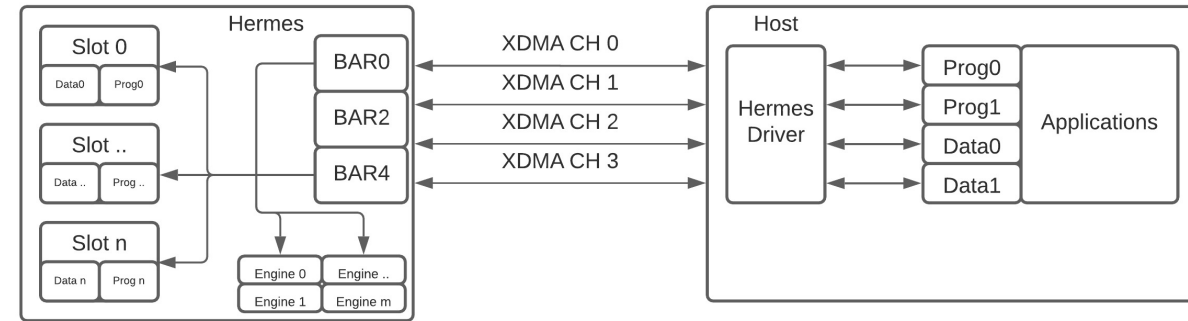
SSD

NVMe:  
NVM/ZNS namespaces  
admin, data, CS command sets  
CS: load/unload/execute

# Mechanisms

- uBPF
  - Added support for registered functions in VM.
- Hermes
  - A BPF accelerator built for AWS F1 instances.
  - Hermes supports loading, unloading and executing programs.
  - Device information and data is stored on BARs (Base Address Registers)
    - The host discover the capabilities of the device from the BAR0.
    - The host uses a driver to offload a BPF program to the BAR2, and then stores it on BAR4.
    - The host executes the BPF program by writing the execute command to BAR0.
      - The device will do an interrupt once the execution is done.
    - The host initiates a DMA transfer to transfer the resulting data back to the host
  - However, Hermes does not support external/registered functions.
    - Thus you can only run what is embedded in the BPF ELF binary itself.
    - The use-cases are therefore limited to very simple programs.

## Hermes





# Mechanisms


- xNVME
  - Our prototype loader will use xNVME to access SSD from the CSD controller.
    - This will make it easy to read from many different storage systems using a single unified interface.
  - Furthermore, we are working on introducing BPF to xNVME, such that xNVME can be used to offload BPF programs to Hermes and our prototype.
    - The interface will be like Hermes: load/unload/execute.
- Open issues relates to exposing the registered functions against which BPF code should be written.
  - **The device may not be attached at compile-time (!)**

# Use-cases



- DAPHNE is an infrastructure for Integrated Data Analysis Pipelines for Large-Scale Data Management, HPC, and Machine Learning funded by EU.
- We are starting with reading CSV.
  - Baseline on host on top of file system
  - First version of BPF-based code, where file access is offloaded
    - Naming issues!
  - Next: offloading parsing and structuring of the data.
    - Many issues (see next slides)

# Use-cases



# DAPHNE

- Finding a balance is the main challenge.
  - DAPHNE is written in C++, which is not easily reducible to BPF code.
    - DAPHNE also uses templating, which cannot be reduced.
  - DAPHNE data structures are scattered in memory, so we'll have to move data to the data buffer before transferring it back to the host.
  - Ideally, we should support the the data structures, but this may be unreasonable due to the number of different data structures in applications.
- Potential solution is reducing data types to bit-lengths and let the host cast the data to the right types.
  - It also introduces the design question on whether the resulting data should be of a structure that the application already implements or a buffer that the application should morph/convert.

DAPHNEs perspective:

long	double	uint
------	--------	------

CSD perspective:

64b	64b	32b
-----	-----	-----

# Conclusion

- Implementation goals stemming from Daphne
  1. push down CSV parsing and structuring via BPF.
    - Issues includes finding the balance in DAPHNE.
      - Create simplified data structures to push to CSD – in C, not C++.
      - Define a structure for the data buffer, that both the runtime and CSD conforms to.
  2. Run on Daisy OpenSSD as CSD
    - This allows us to experiment and collect performance metrics.
    - The open issues include standardizing the interconnect (how data and programs are transferred), implementing xNVME in the prototype and implementing the BPF interface in xNVMe.
- Paper on BPF for computational storage: A Reality Check with DASYS team
  - Analysis of opportunities, challenges and roadblocks.

# Thanks! Questions?

Niclas Hedam | [nhed@itu.dk](mailto:nhed@itu.dk)



IT UNIVERSITY OF COPENHAGEN